

CLAUDE.md Improvement Summary

Comparison of the old CLAUDE (2).md and the new CLAUDE.md

Executive conclusion

The new CLAUDE.md is a clear improvement. It turns the file from a general project description into a stronger operating specification for Claude Code. The new version is clearer about the decision being supported, separates fixed business rules from flexible agent choices, corrects the watch-list benchmark logic, and documents why the correction matters. The file is longer, but only modestly, so it should not create a meaningful context shortage by itself.

What improved

Area	Old version	New version	Why it matters
Structure and decision focus	User, decision, scope, and rules were described, but some parts were blended together.	Adds clearer numbered sections, including User & Decision, Business Rules, and Agent Behaviors.	Claude gets a cleaner operating contract: what the dashboard is for, who uses it, and which choices are non-negotiable.
Business rules	Important rules were spread across multiple sections.	Creates a dedicated Business Rules section.	Reduces the chance that Claude misses rules such as FPA-only scope, no cost figures, no forecasts, and required N counts.
Watch-list benchmark	Recent 3-year performance was compared with a 2006-2025 median annual baseline.	Corrects the benchmark to the 20-year mean annual value, separately for fire count and total hectares.	This is the biggest technical improvement. Hectares burned is highly skewed, so a median denominator can create exaggerated ratios.
Cause definition	Human-vs-Lightning was mentioned but less explicitly defined.	Defines human-related causes as every GENERAL_CAUSE except Lightning, Undetermined, Under Investigation, and Restart.	Makes the cause-mix logic reproducible and prevents inconsistent dashboard calculations.
Iteration evidence	The changelog captured two early mock gaps.	Adds the third gap: the median baseline problem for hectares.	Shows that the mock exposed a real statistical flaw and that the fix was folded back into the spec before build.

Most important improvement

The key upgrade is the correction from a median baseline to a mean annual baseline for the watch-list signal. Median is still appropriate for describing typical individual fire size, because individual fire sizes are extremely right-skewed. However, the watch-list asks whether a region is above its normal annual total. For that purpose, using median annual hectares as the denominator can exaggerate recent ratios when the median year burned very little area. The new version correctly uses the 20-year mean annual value for the watch-list while still reporting fire size as median, 90th percentile, and total hectares.

Length and context impact

The new file is not too long for its purpose. The old version is about 1,678 words, while the new version is about 1,951 words. The increase is roughly 273 words, or around 400 additional tokens. This is small compared with the benefit of preventing repeated scope and calculation errors.

The larger context risks are more likely to come from long CSV previews, pasted logs, full data dictionaries, large code files, and repeated debugging history. The CLAUDE.md file itself should not cause context shortage unless several other long project documents are loaded at the same time.

Recommendation

Keep the new CLAUDE.md as the main operating spec. If the team wants to make it even sharper, move Appendix A and older changelog notes into a separate project-rationale file. Keep the active objective, decision rules, business

rules, agent behaviors, and the median-vs-mean correction inside CLAUDE.md because those directly affect future Claude Code outputs.